

# OS 2017

Homework4: memory allocator implementation

(Due date 01/18 23:59:59)

# Objective

- Understand how to manage heap
- Understand how malloc() and free() work

# Requirements (1/2)

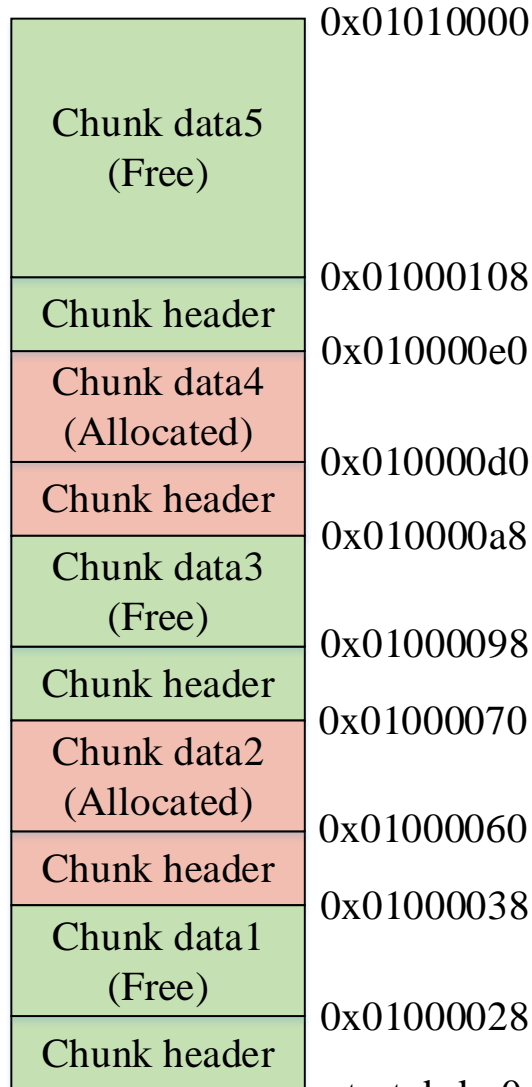
1. Implement a memory allocator library for user application
  - The library must provide the following 3 functions as the interface
    - `void *hw_malloc(size_t bytes);`
      - *bytes*: the required memory size in bytes
      - Return the valid virtual address (starting address of the *data* part) if success; Otherwise, return NULL.
    - `int hw_free(void *mem);`
      - free the virtual memory
      - *mem*: starting address of the *data* part
      - Success: return 1; Fail: return 0.
    - `void *hw_get_start_brk();`
      - Return the starting address of the heap
  - Use chunk (*slides 5-7*), bin (*slides 8-9*) and `sbrk()` to manage heap

# Requirements (2/2)

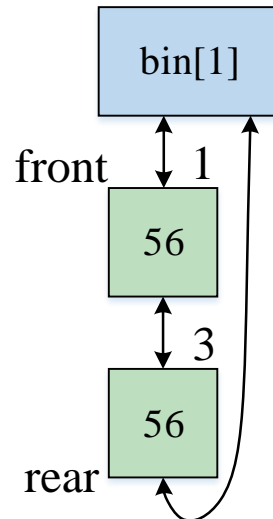
## 2. Write a user application to test the memory allocator library

- Should receive 3 kinds of commands
  - 1) **alloc** *N*
    - Call `hw_malloc(N)` to allocate *N* bytes of **data** memory
    - Print **relative data address** (i.e., offset between `start_brk` and the address returned by `hw_malloc()`)
  - 2) **free** *ADDR*
    - Call `hw_free()` to free the memory at (`start_brk + ADDR`)
    - Print either “success” or “fail”
  - 3) **print** *BIN*
    - Print address and size information of a given bin
    - *BIN* can be `bin[0]`, `bin[1]`, `bin[2]`, `bin[3]`, `bin[4]`, `bin[5]`, or `bin[6]`
- Continuously receive commands from stdin until *EOF* (*Ctrl+D*)
  - Should successfully run  
“`cat testfile.txt | hw4_mm_test > outputfile.txt`”
  - Command input/output format is shown in *slide 4*

# Input/output format



- alloc
  - print (**relative**) **data** address in a line
- free
  - print success/fail in a line
- print
  - print (**relative**) **chunk** address and **size** of each free chunk in the given bin, from the front to the rear
  - print a line for each chunk; pad 8 dash signs between the address and size



## Input example

```

alloc 16
alloc 16
alloc 16
alloc 16
free 0x00000028
free 0x00000098
print bin[1]
    
```

## Output example

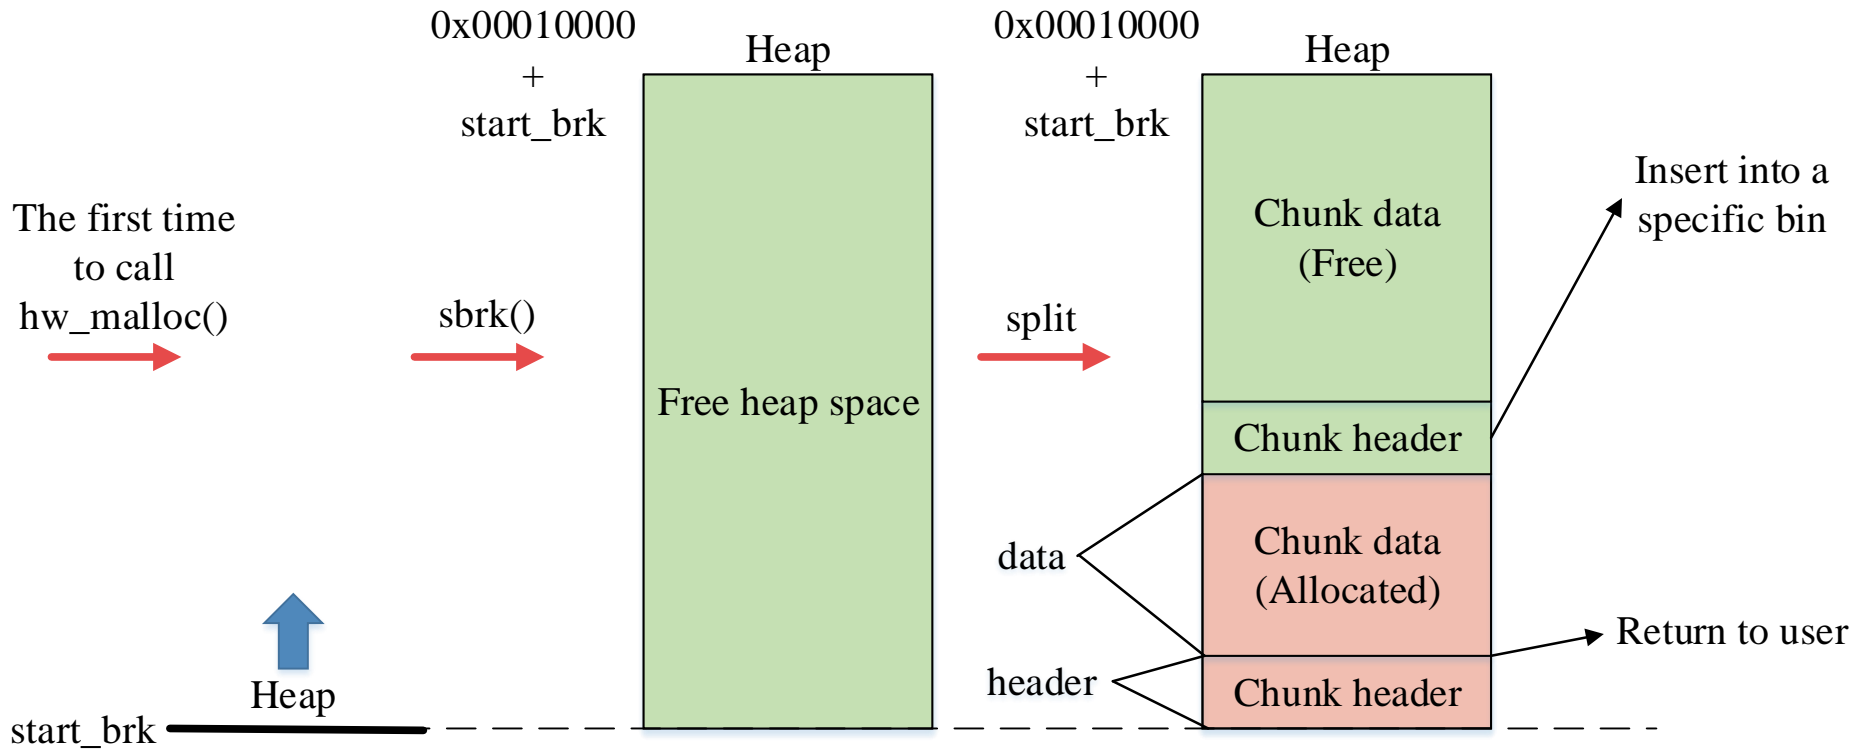
```

0x00000028
0x00000060
0x00000098
0x000000d0
success
success
0x00000000-----56
0x00000070-----56
    
```

# Chunk (1/3)

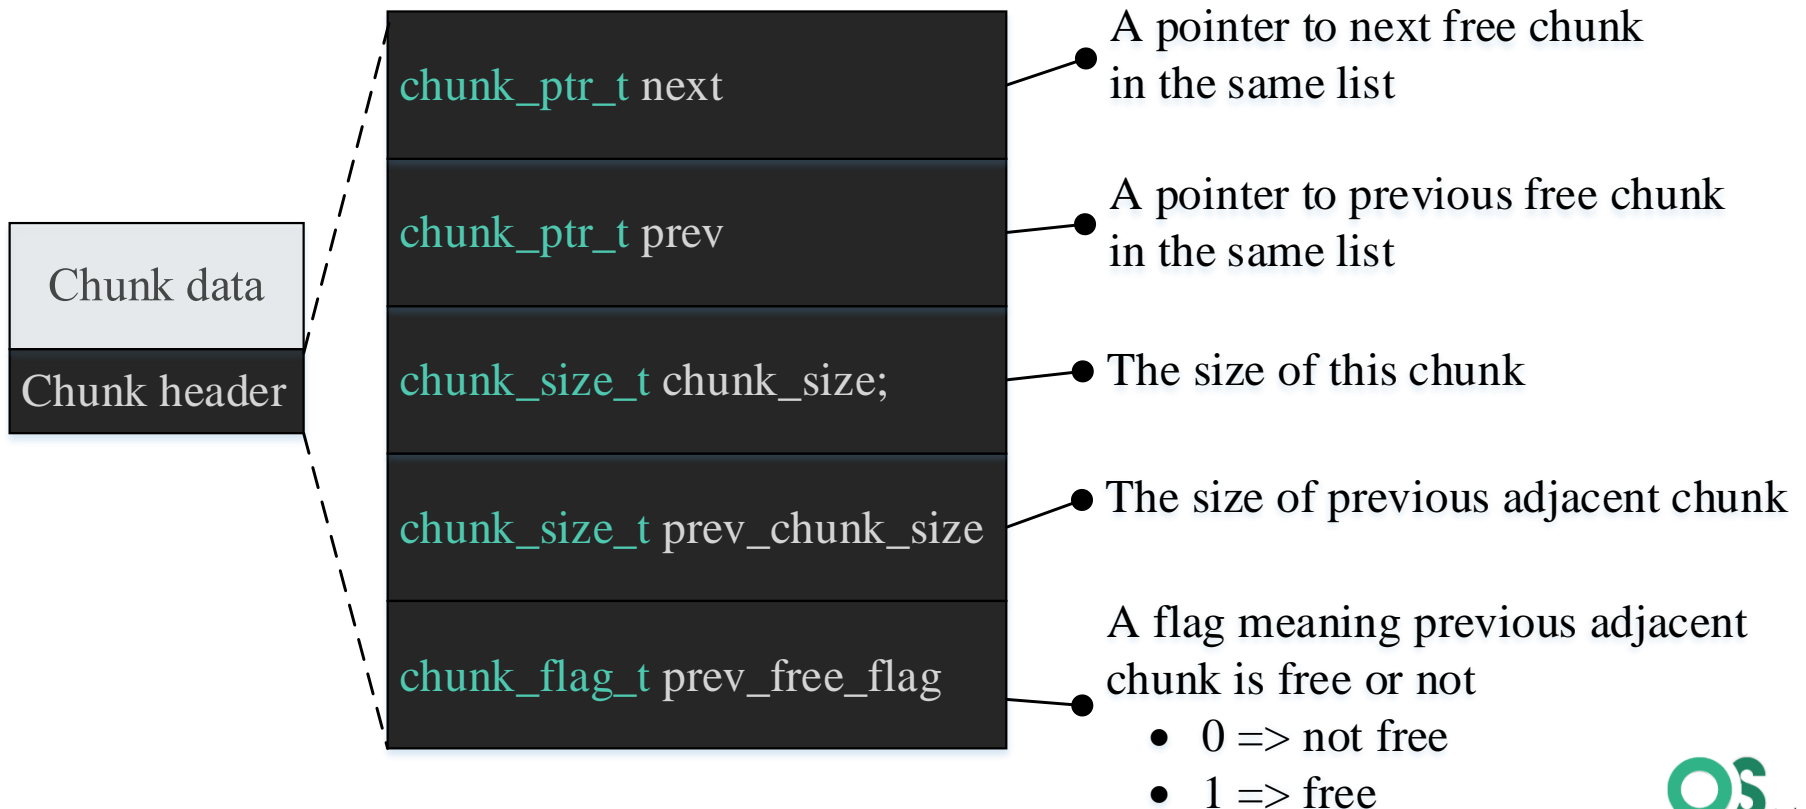
- The continuous heap space is split into chunk(s) for management
  - Chunk size should be multiple of 8 bytes
- When `hw_malloc()` is called for the first time, use `sbrk()` to allocate a **64KB** heap, and then split it into two chunks (*as in slides 6 and 10*)
  - an allocated chunk (**lower** address), returned to the caller
  - a free chunk (**higher** address), inserted in a specific bin
- Each chunk contains two parts, header and data
  - Header (lower address) (*as in slide 7*)
  - Data (higher address), the actual memory space return to caller
- Each free chunk resides in a specific bin (described in *slides 8-9*)
- Adjacent free chunks must be merged (described in *slide 11*)

# Chunk (2/3)



# Chunk (3/3)

- Chunk header (40 bytes)
  - There are 5 members in the header
    - `chunk_ptr_t`, `chunk_size_t`, and `chunk_flag_t` can be defined by yourself, but each of them should be 8 bytes

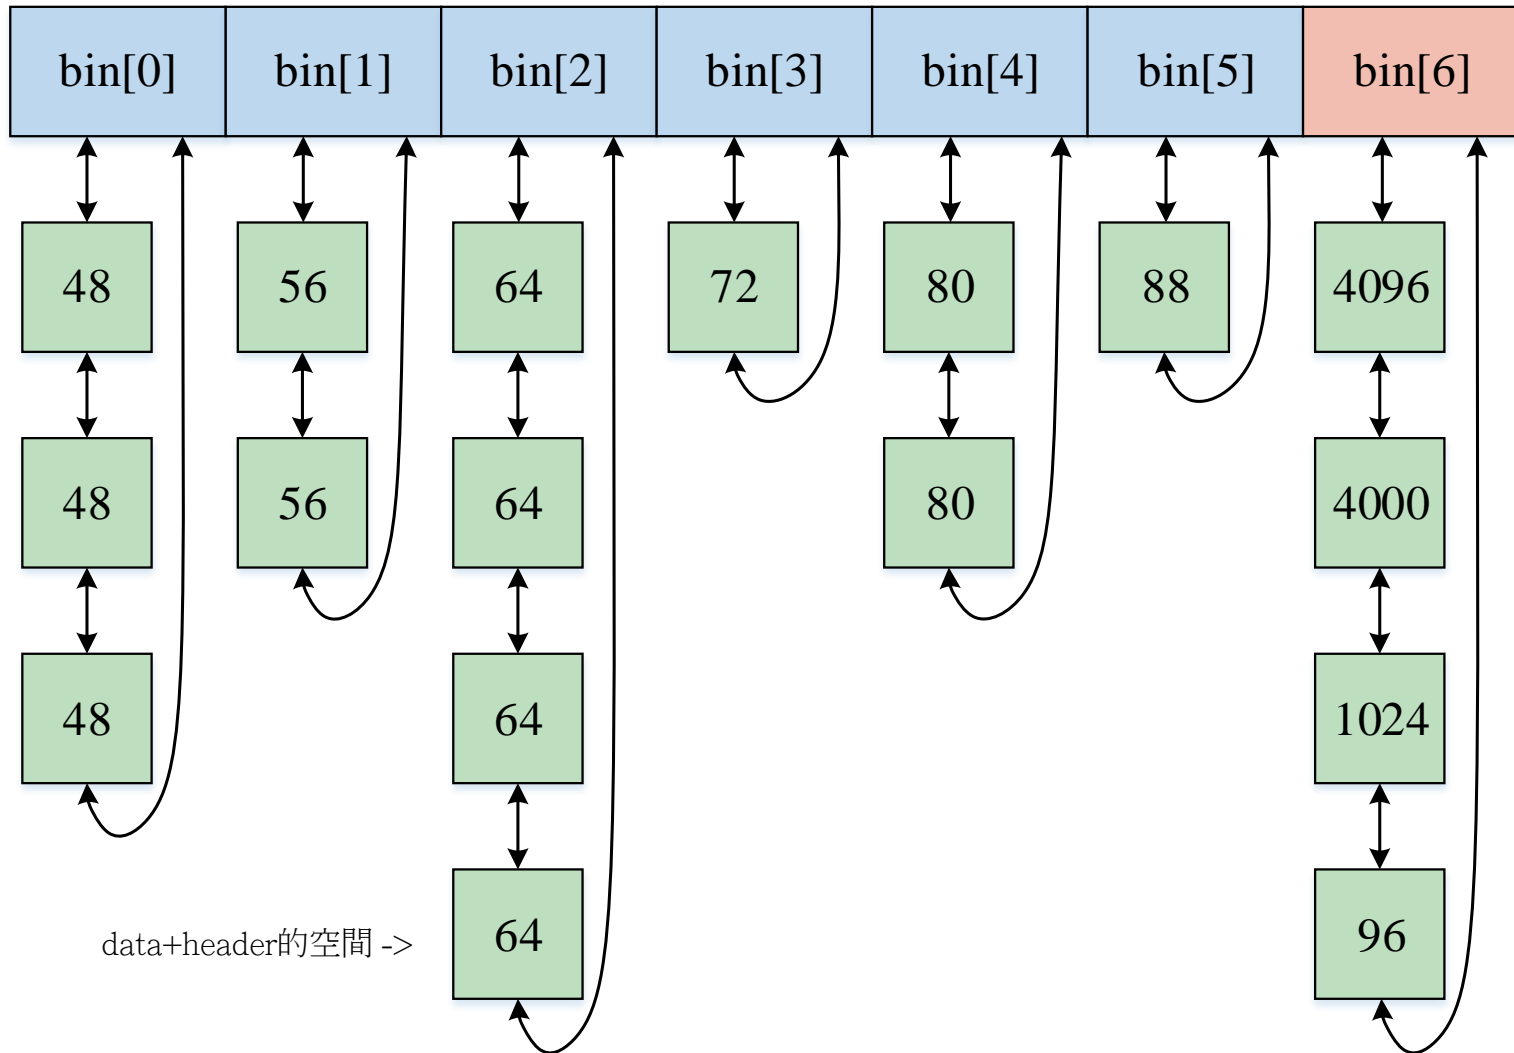




# Bins

- A bin is a circular doubly-linked list of **free** chunk(s) (*slide 9*)
- You should manage 7 bins
  - bin[0]-bin[5] hold chunks with fixed size (*as in slide 9*)
  - bin[6] hold chunks with sizes  $> 88$  bytes
    - Chunks in this bin is sorted by chunk size (in **descending** order)
- Use **best fit** to select a chunk during memory allocation
  - If there are multiple chunks with the same size, select the one that is the nearest to the front(**insert at the rear**) 從前面拿，從後面放
  - You may need to split the selected chunk (*slide 10*)

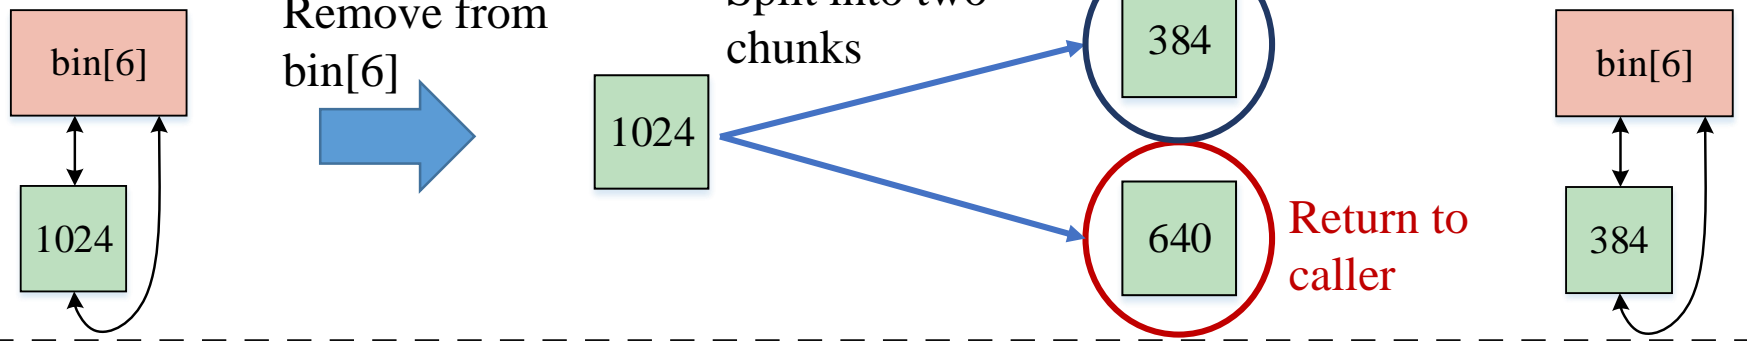
# Bin example




# Split

- When `hw_malloc()` is called, *split* may be performed on the chunk you wish to return to the caller; If split occurs, return the chunk with the **lower** address
- A chunk in `bin[0]-bin[5]` **CANNOT** be split
- A chunk **CANNOT** be split if the remaining size **after split**  $< 48$  bytes (40+8)
- A chunk in `bin[6]` may be split
  - e.g., allocate a 640-byte chunk (40-byte header + 600-byte data) in the following 2 cases

Case1: split



Case2: not split


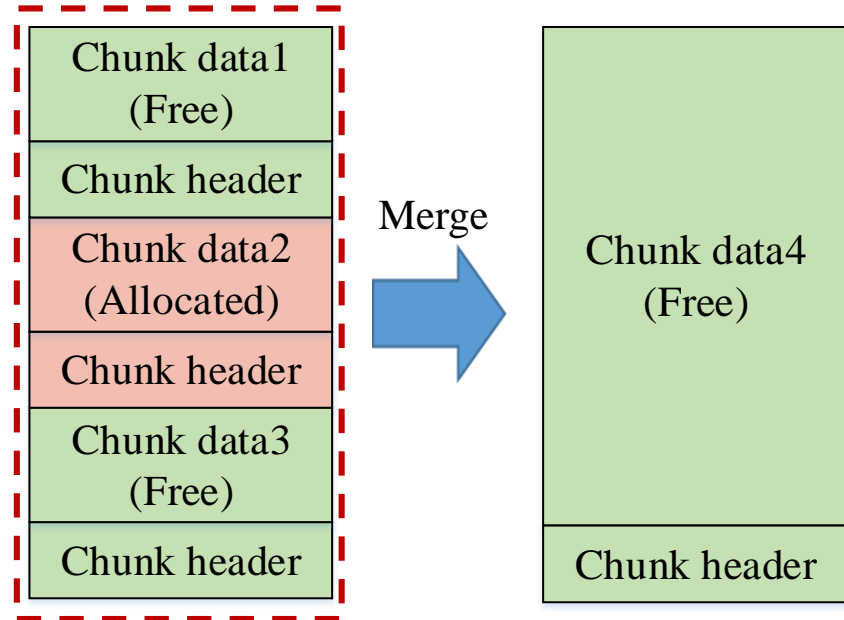


# Merge

- When `hw_free()` is called, adjacent free chunks must be merged into one

For example,

- Chunk2 (48 bytes) is going to be freed
- Chunk1 and Chunk3 are both free and adjacent to Chunk2.
- Chunk1, Chunk2, Chunk3 should be merged (become Chunk4)



# Bonus

- Implement a dynamically growing/shrinking memory allocator library 可以call第二次sbrk長第二個heap，或拿掉heap
- Implement a memory allocator library for multi-threaded process

# References

- sbrk()
  - [Linux man page](#)
- Streams, pipes, and redirects
  - [IBM](#)